

DevOps, Software Evolution and Software Maintenance

Group D

Group Members:

Victor Memborg-Heinrichsen	vmem@itu.dk
Peter Bjørholm Hansen	pbjh@itu.dk
Lukas Vranic	luvr@itu.dk
Anne-Marie Rommerdahl	annro@itu.dk

BSc Software Development
IT University of Copenhagen
BSDSESM1KU
May 27, 2025

Contents

1	System's perspective	1
1.1	Architecture of MiniTwit	1
1.2	Dependencies of MiniTwit	2
1.3	System Interactions	2
1.4	Current state of MiniTwit	2
2	Process' perspective	2
2.1	Security Assessment	2
2.2	Scaling and Upgrades	5
2.3	CI/CD Chain	5
2.4	Monitoring	9
2.5	Logging	10
3	Reflection Perspective	10
3.1	Evolution and Refactoring	10
3.2	Operation	10
3.3	Maintenance	11
4	Usage of AI-assistants	11

1 System's perspective

Our MiniTwitter is an 'X' (formerly known as Twitter) clone written in Golang using the Gorilla web framework. It is a project that continues development on the ITU-MiniTwitter system presented in the course DevOps, and it consists of a web service and an API service.

1.1 Architecture of MiniTwitter

MiniTwitter follows the server-client architecture, and the server is deployed via a Docker Swarm, which consists of virtual machines (aka 'droplets') on DigitalOcean. Data, such as user-data, messages and followers are stored in a PostgreSQL database, which is also hosted on DigitalOcean.

In the swarm, we have manager-nodes and worker-nodes. The only services allowed to run on the managers are Prometheus and Dozzle (more about the swarm services in section 4). Figure 1 is a Specification Level Deployment Diagram, which shows how ITU-MiniTwitter is deployed on a worker-node.

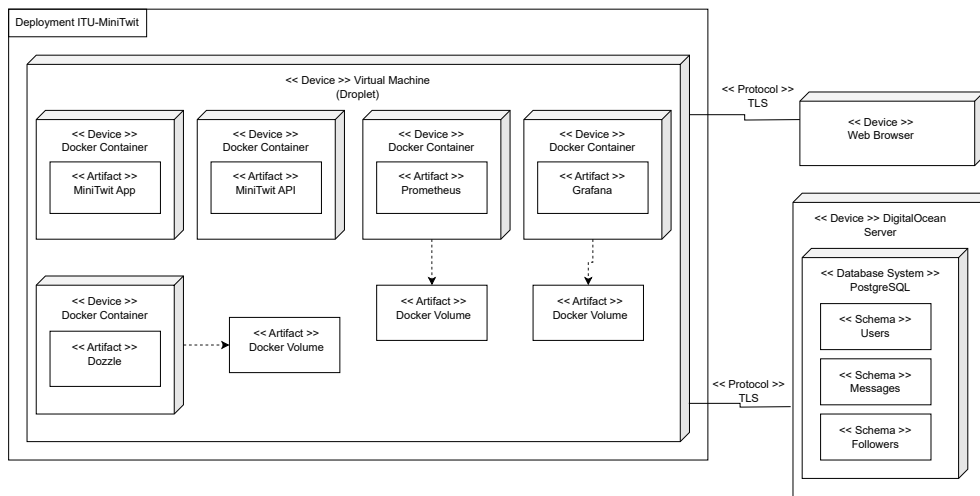


Figure 1: WIIP

1.2 Dependencies of MiniTwit

1.3 System Interactions

1.4 Current state of MiniTwit

We have implemented the two tools Sonarqube and Code Climate, which can assist in estimating the maintainability and technical debt of our project.

Code Climate

Code Climate reports 10 code smells, 0 duplications and 0 other issues. This leaves us with an A-rank (see figure 2)

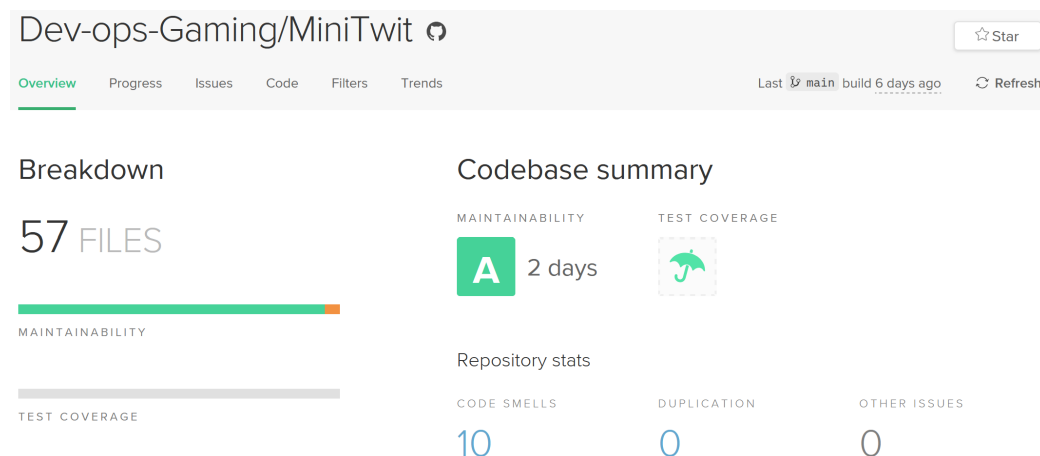


Figure 2: Bitch?

Sonarqube

Sonarqube reports 18 issues - all in the category 'maintainability' - and 0.9% code duplication (see figure 3). A majority of these issues are found in the tests, and many of the issues are in regards to string duplicates.

2 Process' perspective

2.1 Security Assessment

To protect our system against adversaries, we first identify the assets which must be protected. We then identify which threats we could be faced with and how.

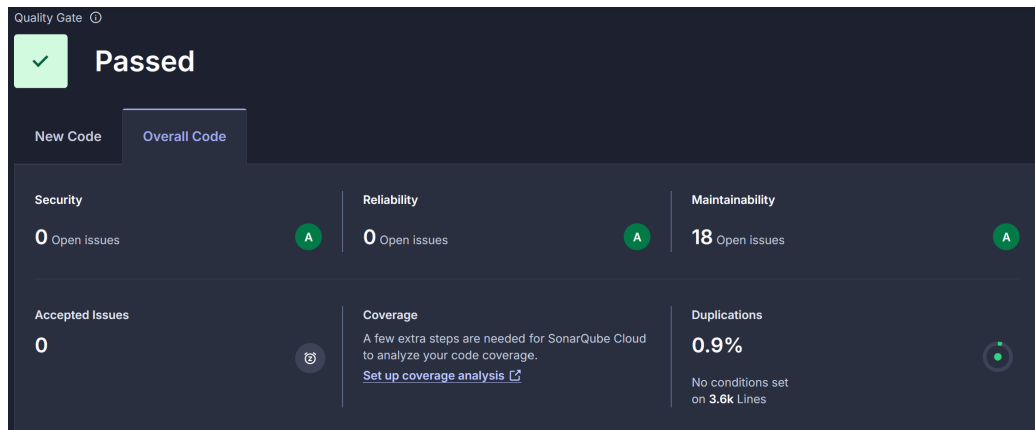


Figure 3: We need to fix position of this when report almost done

Finally, we analyze these scenarios to figure out which scenarios pose the biggest threats.

(i dunno if this intro text is necessary. Can maybe be cut if we don't have space?)

Risk Identification

Our assets:

- Codebase
- Database
- Logs

Threat Sources:

- SQL Injection
- Exposed ports (Prometheus)
- Vulnerable password hashing
- No requirements for user password (length, symbols, etc.)
- Vulnerable libraries
- No password on Dozzle (logging)

Risk Scenarios:

- A - Attacker performs SQL injection to download sensitive user data
- B - Attacker performs SQL injection to delete information from database
- C - Attacker performs SQL injection to log in on a legitimate user's account
- D - Attacker brute-forces the hash of a password, letting them log in on a legitimate user's account

- E - Attacker uses the open ports to overload our Prometheus with queries
- F - Attacker intercepts messages sent over the network, as they are not encrypted
- G - Attacker exploits vulnerable library, which lets them perform some malicious attack
- H - Attacker accesses our logs on Dozzle, which lets them see all outputs, like usernames

Risk Analysis

Here are the scenarios presented on the risk matrix:

	Rare	Unlikely	Likely	Certain
Catastrophic	B			G
Critical	A			G
Marginal			F	E,G
Negligible	C		D	G,H

Scenario A, B and C have a 'Rare' probability, as we already utilize the ORM-library GORM. GORM uses prepared statements and automatically escapes arguments to avoid SQL injection. Without GORM the probability would be much higher.

Scenario F was dealt with using the TLS protocol. Our Minitwit application using TLS is hosted at <https://lukv.dk>.

For scenario E, we would have to perform a workaround, as the exposed ports are caused by Docker and UFW being incompatible, an issue which has not yet been patched. However, Prometheus is the only service which is exposed, and the only threat is Prometheus crashing due to an overload of queries. Therefore, it is not considered a high-priority threat.

We could reduce the chance of scenario G happening, by keeping our libraries updated. We could also use tools that detect vulnerabilities in dependencies, and keep the dependencies updated automatically. Scenario H can be prevented by locking our Dozzle interface with a password. However, no sensitive data is logged in Dozzle, and thus it is considered low priority.

2.2 Scaling and Upgrades

Scaling

To scale our Minitwit application, we performed horizontal scaling using Docker Swarm. We created three manager nodes and four worker nodes.

We have replicas of the following services:

Service	Replicated/Global	No. of replicas
Minitwit App	Replicated	3
Minitwit API	Replicated	4
Prometheus	Global	N/A
Dozzle	Global	N/A

Prometheus and Dozzle are set to be global services, meaning one instance of each runs on every node in the swarm. This ensures that monitoring and logging are consistently performed on all nodes, preventing any loss of critical information, whether it is logging or monitoring, due to missing coverage.

Upgrades

We use a rolling update strategy with a start-first order to update the services in our swarm. This means one replica at a time is updated by starting a new container before stopping the old one, ensuring minimal downtime. Updates are monitored for 30 seconds, and if any failures are detected, our swarm will automatically roll back to the previous working version [1].

2.3 CI/CD Chain

We are using GitHub Actions to automate the testing and deployment of Minitwit. There are 5 GitHub Actions workflows:

1. Deploy services to DigitalOcean
2. Use Hadolint on dockerfiles
3. Run golangci-lint tool
4. Release MiniTwit (automatically)

Use Hadolint on dockerfiles

This workflow only runs when it get called by other workflows, and analyse if the dockerfiles `docker/Dockerfile` and `docker/Dockerfile.api` contains any linting issues.

The steps for this workflow is:

1. Checkout

(a) Purpose

- i. The workflow uses `actions/checkout@v2` to fetch the code from the repository.

(b) Steps

- i. Checkouts the code.

2. Hadolint Action app

(a) Purpose

- i. Checks for any linting errors using `hadolint/hadolint-action@v3.1.0` in `docker/Dockerfile`.

(b) Steps

- i. Runs `hadolint-action` on `docker/Dockerfile`.

3. Hadolint Action api

(a) Purpose

- i. Checks for any linting errors using `hadolint/hadolint-action@v3.1.0` in `docker/Dockerfile.api`.

(b) Steps

- i. Runs `hadolint-action` on `docker/Dockerfile.api`.

Run golangci-lint tool

This workflow also only runs when it get called by other workflows, and analyses the go source code any linting issues.

1. Environment setup

(a) Purpose

- i. Defines some environment variables that are used in the workflow.

(b) **Steps**

- i. `GO_VERSION = stable`
- ii. `GOLANGCI_LINT_VERSION = v1.64`

2. Detect Modules

(a) **Purpose**

- i. To output all the Go modules in `./minitwit`.

(b) **Steps**

- i. Checkout the repository.
- ii. Runs `go list -m` to list all Go modules and outputs in JSON format.

3. Format Go Files

(a) **Purpose**

- i. To format all the `.go` files, done by using `Jerome1337/gofmt-action@v1.0.5`.

(b) **Steps**

- i. Checkout the repository.
- ii. Runs `gofmt-action` to verify formatting in `./minitwit`.

4. Golangci Lint

(a) **Dependency**

- i. Needs *Detect Modules* to run.

(b) **Steps**

- i. Checkouts the repository and setups a go environment.
- ii. Runs `golangci-lint` in each GO module.

Deploy services to DigitalOcean

This workflow runs every time the main branch gets a push.

1. call-hadolint and call-golangci

(a) **Dependency**

- i. *Hadolint on dockerfiles* and *golangci-lint*.

(b) **Purpose**

- i. to ensure code quality before deployment.

(c) **Steps**

- i. Checkout the repository.
- ii. Run `Hadolint` on `dockerfiles`.
- iii. Run `golangci-lint`.

2. **Run tests**

(a) **Dependency**

- i. `call-hadolint` and `call-golangci`.

(b) **Purpose**

- i. To test the code to ensure code quality before deployment.

(c) **Steps**

- i. Checkout the repository.
- ii. Setup a Go environment with v1.22 of Go.
- iii. Setup docker Compose.
- iv. Find the test script (`run_tests.sh`).
- v. Run the test script.

3. **Build And Deploy**

(a) **Dependency**

- i. *Run tests*

(b) **Purpose**

- i. Build a Docker image, push it to Docker Hub and Deploy it on Digital Ocean Droplet.

(c) **Steps**

- i. Checkout the repository.
- ii. Verify secrets.
- iii. Login to Docker Hub.
- iv. Set up Docker Buildx.
- v. Build and push `minitwit-app`.
- vi. Build and push `minitwit-api`.
- vii. Configure SSH.
- viii. Sync files with `rsync`.
- ix. Deploy to server.

Release MiniTwit (automatically)

This workflow creates a release on Github every thursday at 23:30 UTC.

1. Automatic Release

(a) Purpose

- i. Creates a release on Github.

(b) Steps

- i. Checkout the repository.
- ii. Fetch latest version tag.
- iii. Determine next version.
- iv. Validate against SemVar pattern.
- v. Generate release notes.
- vi. Create GitHub Release.

2.4 Monitoring

We monitored MiniTwit using Prometheus for collecting metrics and Grafana for visualizing them. This setup provided great insights into the application's performance and behavior.

Prometheus

Prometheus was integrated into both the API and the application through a custom middleware that intercepts HTTP requests to gather metrics. The metrics collected includes:

- `http_requests_total`: A counter that counts the total number of HTTP requests. It includes labels for request path, HTTP method, and status code. This metric helped us monitor request load per endpoint, track response statuses (in general or just on an endpoint basis).
- `http_request_duration_seconds`: A histogram measuring the time taken to process HTTP requests, in seconds. It is labeled by request path and method, enabling performance benchmarks for endpoints.
- `http_response_messages_total`: A counter that logs the total number of HTTP responses, categorized by status code and message type. Before logging was fully implemented, this metric was particularly helpful in identifying which endpoints triggered specific status messages and understanding the reasons behind them. A bit deprecated as soon as logging was implemented.

Seen in retrospect, it would have been a good idea also to collect metrics on database queries, but this will be described more thoroughly in section 3.3

Grafana

While Prometheus was important for collecting data, Grafana is utilized for visualizing the data. Grafana connects to Prometheus as a data source, enabling the creation of dashboards. As mentioned in section 2.4 we monitored both our API and app, which resulted in us setting two dashboard up; one for the app and one for the API.

For instance, the API dashboard converted the metrics collected by Prometheus into a clear visual representation of the API's health and performance. Panels were set up to showcase:

- Indicators from `http_requests_total`, such as the rate of requests per second, overall request load distributed by endpoint, and the ratio of successful and client-error status codes.
- Performance monitored by `http_request_duration_seconds`, including average response times and 99th percentile latency, allowing for quick identification of performance degradation.

2.5 Logging

We implemented Logging in MiniTwit using Dozzle. We were recommended to implement an ELK or ELFK stack, but for our group Dozzle was a better lightweight alternative, I will explain more about why this is later on.

3 Reflection Perspective

3.1 Evolution and Refactoring

3.2 Operation

API issue

The very first issue we had was the API becoming unavailable due to crashing when the database encountered an error. We had a hard time debugging this, as the crash did not provide any valuable information in the Docker logs (a logging tool had not been introduced yet), only to discover that we used `panic`, which causes the program to terminate in Go. This took a bit of time to realize, but when we finally

traced the issue, we replaced the panic call with returning a proper status code and tried to enforce more thorough code reviews.

Adding indexes

As we monitored the performance of our API and web application throughout the project, we noticed a degradation in response times. At one point, the `/public` endpoint was particularly affected, with average response times reaching up to 70 seconds. This was unacceptable, and we began investigating the cause.

Although we had not implemented database query monitoring through prometheus, we suspected that the bottleneck might lie in our SQL queries. Looking at the performance metrics provided by DigitalOcean, where our PostgreSQL database is hosted, we identified that the `queryMessages` method was the cause of this issue, and it became clear that we needed to index on frequently queried columns. The indexation resolved the issue, and our `/public` endpoint has since then had a response time of 20-50ms. In future projects, it will be optimal to monitor the database execution time and visualize it as a `Time series` in Grafana to see if queries degrade over time.

3.3 Maintenance

4 Usage of AI-assistants

We utilized two large language models (LLMs) throughout the project to support our development process: GitHub Copilot and ChatGPT. GitHub Copilot was primarily used for code completion and was constantly active during development. ChatGPT, on the other hand, was used to quickly gain an overview of how to attack an issue, e.g., the database index issue described in section 3.2 or a tutor conveying the essential documentation for e.g., our web framework "Gorilla".

We found that the LLMs significantly improved our workflow by accelerating learning and implementation. However, we suspect Copilot of introducing the API issue also described in 3.2, so it is important to be critical when utilizing these tools.

References

- [1] Docker, Inc. *Compose Deploy Specification*. n.d. URL: <https://docs.docker.com/reference/compose-file/deploy/> (visited on 05/25/2025).